

---

CSE 3400/CSE 5850 - Introduction to Computer & Network  
Security  
/ Introduction to Cybersecurity

Lecture 12  
Public Key Infrastructure

Ghada Almashaqbeh  
UConn

**\*Adapted from the textbook slides**

---

---

# Outline

- ❑ Motivation.
- ❑ Public key infrastructure (PKI) components.
- ❑ PKI goals.
- ❑ X.509 PKI concepts.
- ❑ Intermediate CAs and trust path verification.
- ❑ Certificate revocation.


---

# Public keys are very useful...

- Secure web connections.
- Software signing
- Secure messaging, email.
- Cryptocurrency and blockchains.
- But ... how do we know the public key of an entity? And how can we trust that this entity is indeed who claims to be and owns a specific public key?
  - Mainly: the key must be signed by a **trusted Certificate Authority (CA)**.
- *Public key infrastructure (PKI)* defines how to issue, manage and use such certificates.

# Public Key Certificates & Authorities

- **The big picture:** when receiving a party's (the **subject**) public key, it will be accompanied with a certificate.
  - A valid certificate means that the entity is who claims to be and she owns the corresponding the public key.
- **Certificate:** signature by a **Certificate Authority (CA)** over subject's public key and attributes.
  - A chain of trust covering multiple intermediate CAs back to the root CA.
- **Attributes:** identity (ID) and others...
  - Validated by CA (liability?)
  - Used by **relying party** for decisions (e.g., use this website?)



---



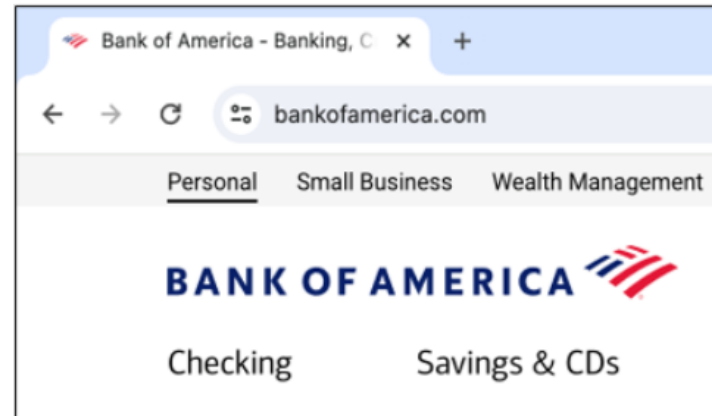

# Certificates are all about **Trust**

- Certificate:  $\psi_{Bob} = \text{Sign}_{CA.s}(\text{Bob.com}, \text{Bob.e}, \dots)$ 
  - CA attests that Bob's public key is *Bob.e*
- Do we **trust** this attestation to be true?
- Special case of **trust management**
  - Important problem far beyond PKI... still not resolved!

# Rogue Certificates

- Rogue certificates: certificates that contain wrong or misleading information.
  - ❑ So they should fail PKI validation.
- Attacker goals:
  - ❑ Impersonate: web-site, phishing email, signed malware..
  - ❑ Equivocating (same name): circumvent name-based security mechanisms, such as *blacklists*, *access-control* ...
- Types of misleading names:
  - ❑ Combo names: bank.com vs. **accts-bank.com**, **bank.accts.com**, ...
  - ❑ Domain-name hacking: accts.bank.com vs. **accts-bank.com**, ... or **accts-bank.co**
  - ❑ Homographic: paypal.com [l is L] vs. **paypal.com** [i is l]
  - ❑ Typo-squatting: bank.com vs. **banc.com**, **baank.com**, **banl.com**,...

# Example of Homographic Attacks



---


# PKI Failures

- Although the signature over the certificate verifies correctly, there is still a failure and the certificate must be revoked.
  - This is called a PKI failure.
- PKI failures include:
  - Corrupt CA.
  - Validation failure.
  - Exposed CA private key.
  - Cryptanalysis certificate forgery.
    - Find collisions in the hash function used in the HtS paradigm,
    - or exploit some vulnerability in the digital signature scheme used for signing.



# Some Infamous PKI Failures

| CA, year(s)          | Description and Reference   |
|----------------------|---|
| Verisign, 2001       | VeriSign issues Microsoft code-signing certificates to attacker [167].  |
| Thawte, 2008         | Validation failures of Thawte and StartSSL [455].   |
| Comodo, 2008         | CertStar, a reseller of Comodo, issued certificates without validation [299].   |
| Comodo, 2011         | Rogue certificates for major sites (e.g., Gmail) [261][291][343].   |
| DigiNotar, 2011      | DigiNotar CA compromised, 531 rogue certificates found, including for *.google.com, used for MitM against Iranian users [291][440].         |
| TurkTrust, 2011-2012 | TurkTrust issued intermediary CAs certificates to end entities; abused to issue certificate for *.google.com (detected on Dec. 2012) [282]. |
| Trustwave, 2012      | Trustwave issued intermediary CA certificate for eavesdropping [368].   |
| ANSSI, 2013          | ANSSI (French CA) issued intermediary CA certificate for MitM [445].  |
| NICCA, 2014          | Intermediary CA NICCA (India) issued rogue certs for Google domains [284].  |
| CNNIC, 2015          | Rogue certificates issued by MCS (Egypt), certified by CNNIC (China) [152][332].  |
| WoSign, 2015         | WoSign and StartCom (owned by WoSign) removed from revoked as CAs after validation and other failures                                       |
| Symantec, 2015-17    | Symantec issued unauthorized certs for over 176 domains [352].  |
| DarkMatter, 2019     | Mozilla, Google revoke intermediary CA of surveillance firm DarkMatter [68], refuse to make it a root CA.                                   |
| Let's Encrypt, 2020  | Let's Encrypt detected a bug in their CAA-validation code, affecting 3 million certificates [1].  |
| TrustCor, 2022       | Root CA TrustCor exposed as related to Spyware [318].   |



# PKI Goals/Requirements



**Trustworthy issuers:** Trust anchor/root CAs and Intermediary CAs; Limitations on Intermediary CAs (e.g., restricted domain names)



**Accountability:** identify issuer of a given certificate



**Timeliness:** limited validity period, timely **revocation**



**Transparency:** public log of all certificates; no 'hidden' certificates!



**Non-Equivocation:** one entity – one certificate



**Privacy:** why should CA know which site I use?

---

# X.509 Certificates

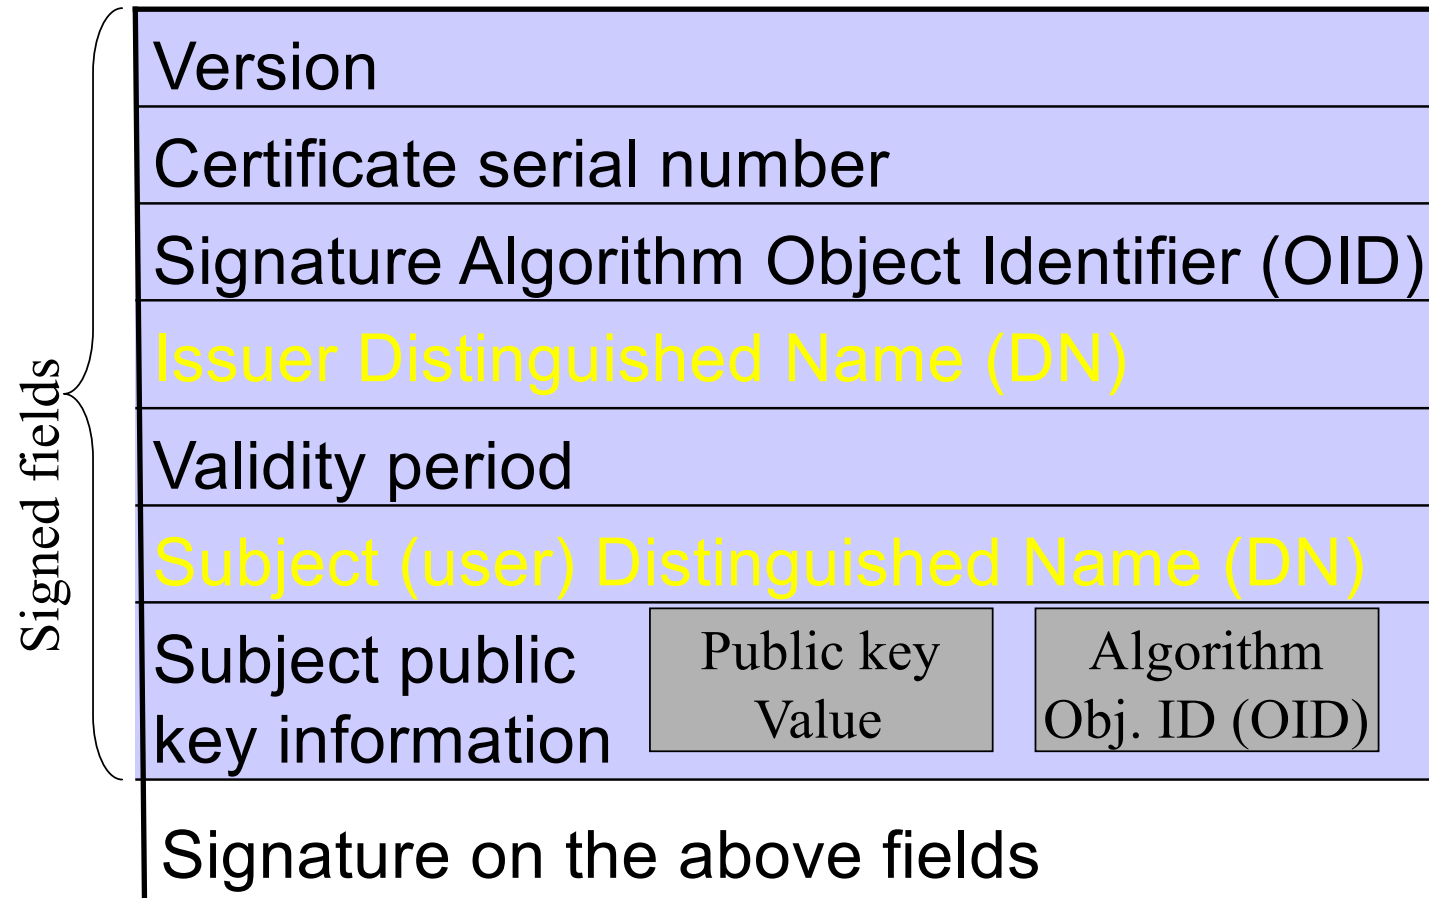
*Part of the X.500 Global Directory Standard*

---

# The X.509 Standard Certificate Format

- Published by ITU (International Telecommunication Union) in 1988 as part of the X.500 global directory standard.
- Idea: Signature binds **public key** to distinguished name (DN) and to other attributes
  - Some defined in X.509 standard, others in `extensions`
- Used widely despite complaints about its complexity.
  - SSL/TLS, code-signing, IP-Sec, ...

# X.509 V1 Certificate Format



# X.509 V1 Certificate Format

- **Version:** the version of X.509 (for V1 it is 1 and so on).
- **Certificate serial number:** a serial number of the certificate, unique among all the certificates issued by this CA.
- **Signature algorithm OID:** an object identifier (OID) for the signature algorithm used to sign the certificate.
- **Issuer DN:** the Distinguished Name (DN) of the issuer of the certificate.
- **Validity period:** the period during which the certificate is supposed to be valid.
- **Subject DN:** the Distinguished Name (DN) of the subject of the certificate, i.e., the entity to whom the certificate was issued.
- **Subject public key information:** includes two parts, one containing the certified public key, and the other providing an OID to identify the public key algorithm with which this public key is to be used.
- **Signature (produced by CA):** a signature over the above fields.

---

# X.509 Certs & Subject Identifiers

- V1: Distinguished Name (for subject & issuer)
- V2: Unique identifiers (for subject & issuer)
- V3: Extensions (used in practice)
  - Some defined in X.509, others elsewhere

# X.509 Certificate Format – Later Versions

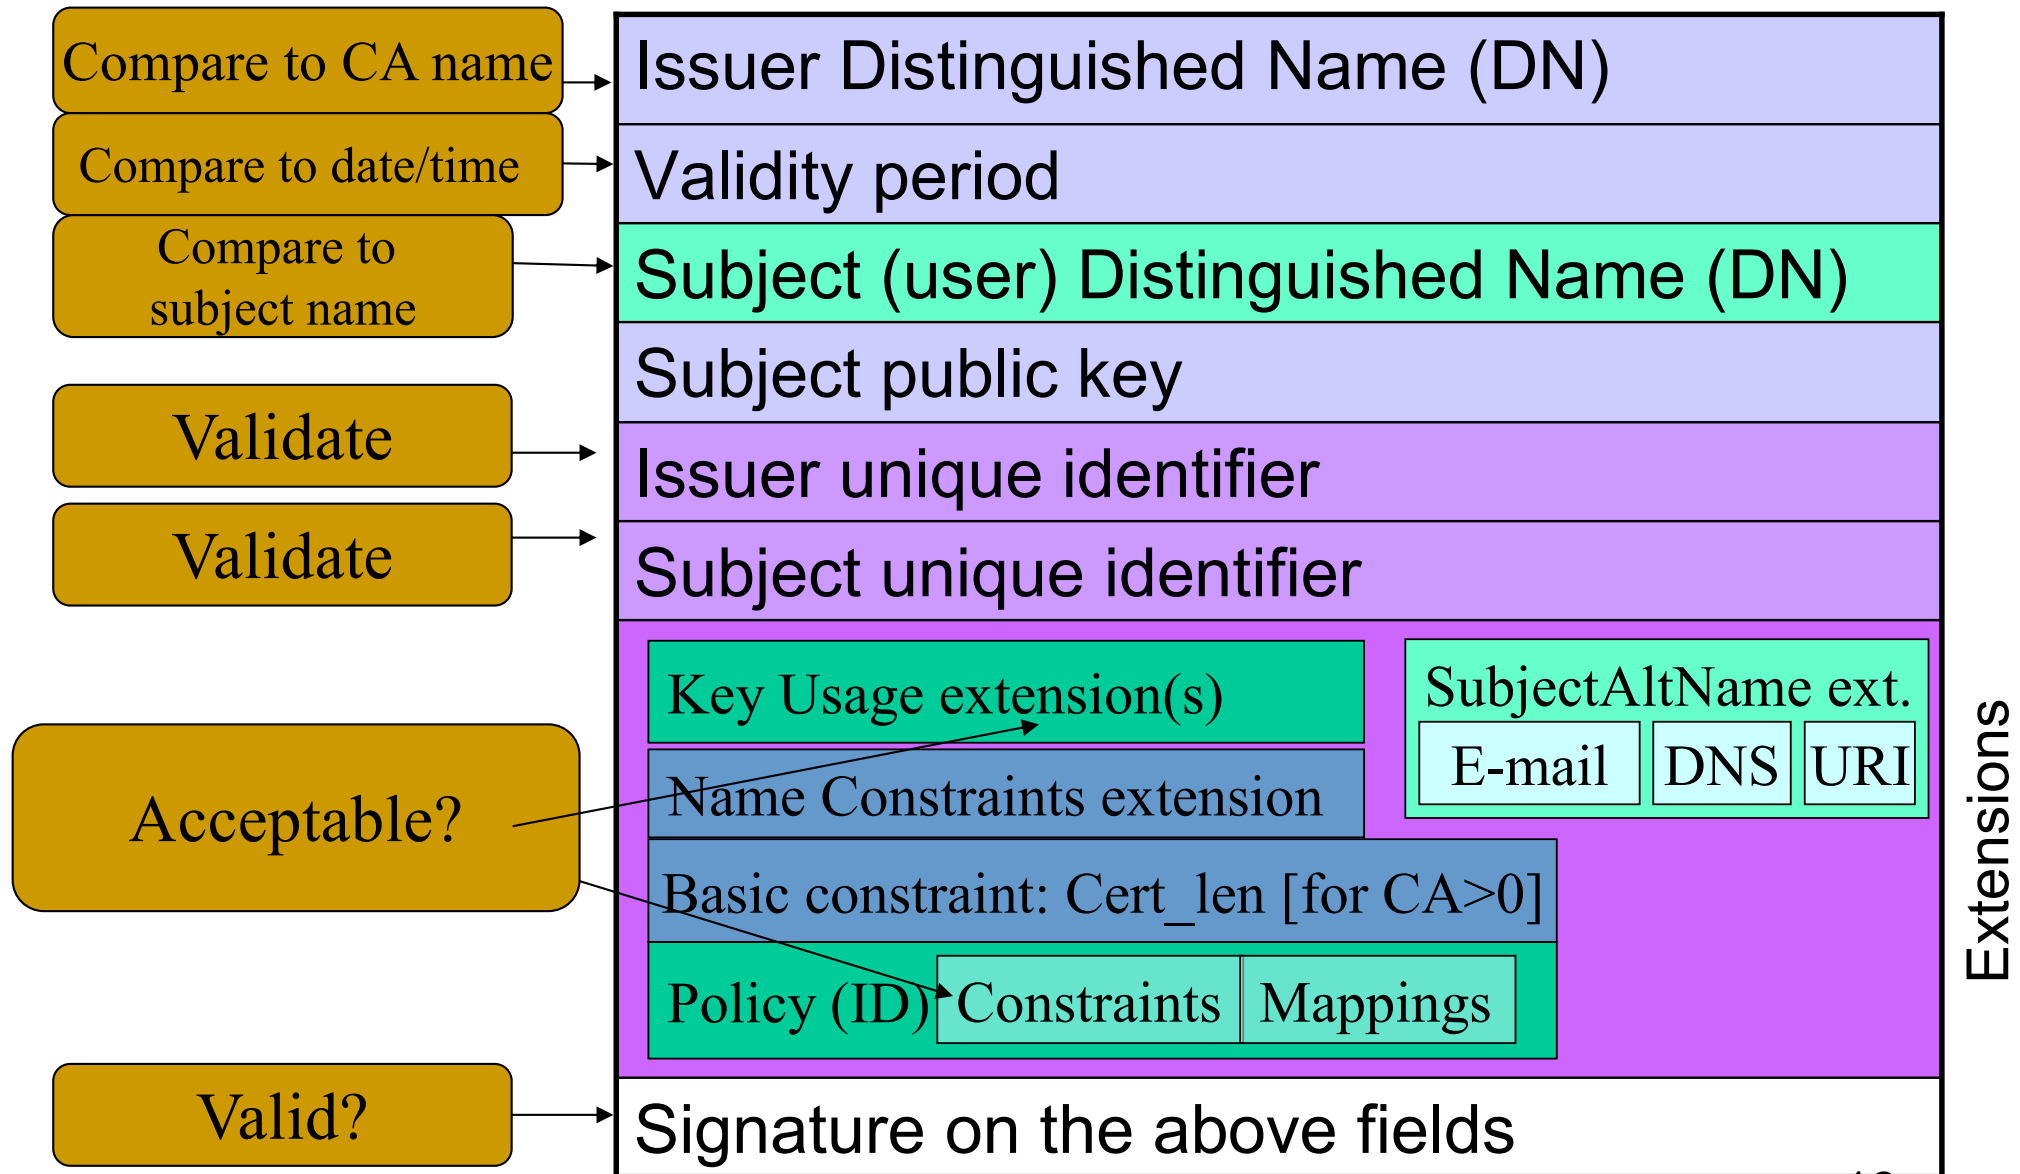
|                               |   |                  |                         |
|-------------------------------|---|------------------|-------------------------|
| Signed fields                 | Version                                     |                  |                         |
|                               | Certificate serial number                   |                  |                         |
|                               | Signature Algorithm Object Identifier (OID) |                  |                         |
|                               | Issuer Distinguished Name (DN)              |                  |                         |
|                               | Validity period                             |                  |                         |
|                               | Subject (user) Distinguished Name (DN)      |                  |                         |
|                               | Subject public key information              | Public key Value | Algorithm Obj. ID (OID) |
|                               | Issuer unique identifier (from version 2)   |                  |                         |
|                               | Subject unique identifier (from version 2)  |                  |                         |
|                               | Extensions (from version 3)                 |                  |                         |
| Signature on the above fields |   |                  |                         |



# X.509 Certificate Format – Later Versions

- **Issuer and subject unique identifiers (V2):**
  - ❑ Added to ensure uniqueness to handle situations where the DN may fail to ensure uniqueness.
  - ❑ Not widely used.
- **Extensions (V3):**
  - ❑ Additional fields to increase the expressiveness of X.509 certificates to facilitate more applications and end users.
  - ❑ Examples include limitations on which application the certificate or public key can be used for, certificate path constraints, policy constraints, etc.
  - ❑ We will not cover these in this course. More details are in a Network Security course.

# X.509 Certificate Validation (simplified)



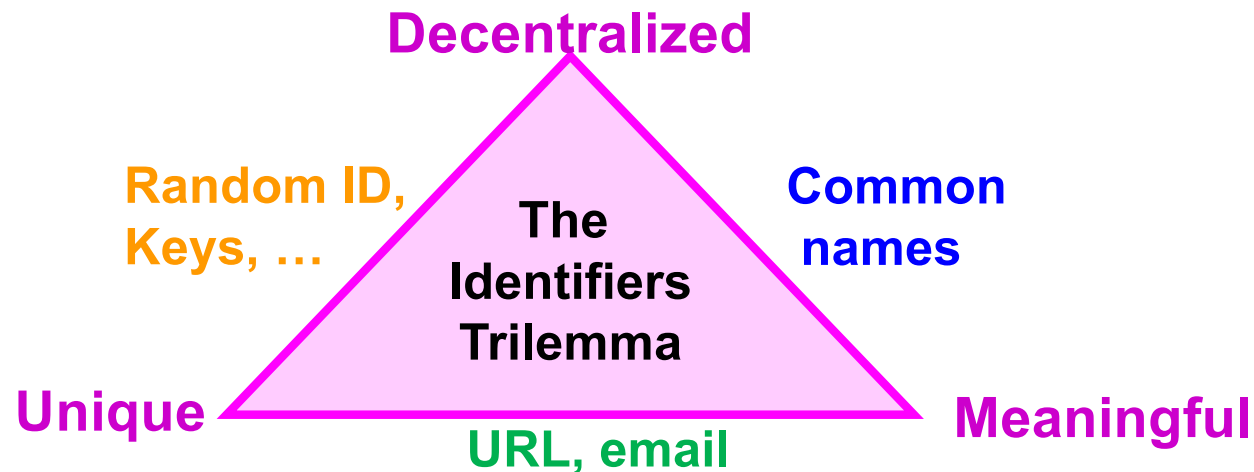
---

# Distinguished Names

- Most certificates contain identifiers.
- Influenced by telecommunication providers.
  - Phone directory services are based on common names.
- Basic goals of identifiers:
  - **Meaningful** (to humans)
    - Memorable, reputation, etc.
  - **Unique** identification of entity (owner)
  - **Decentralized** - with accountability:  
assigned by trusted (certificate) authorities
    - Accountability: identification of the signing authority

# The Identifiers Trilemma

- Achieving the three goals: Meaningful, Unique, Decentralized, seems very challenging!
- Examples of achieving any two of the goals:
  - Unique + Meaningful: URL, email
  - Meaningful + Decentralized: common name
  - Unique + Decentralized: hash of key

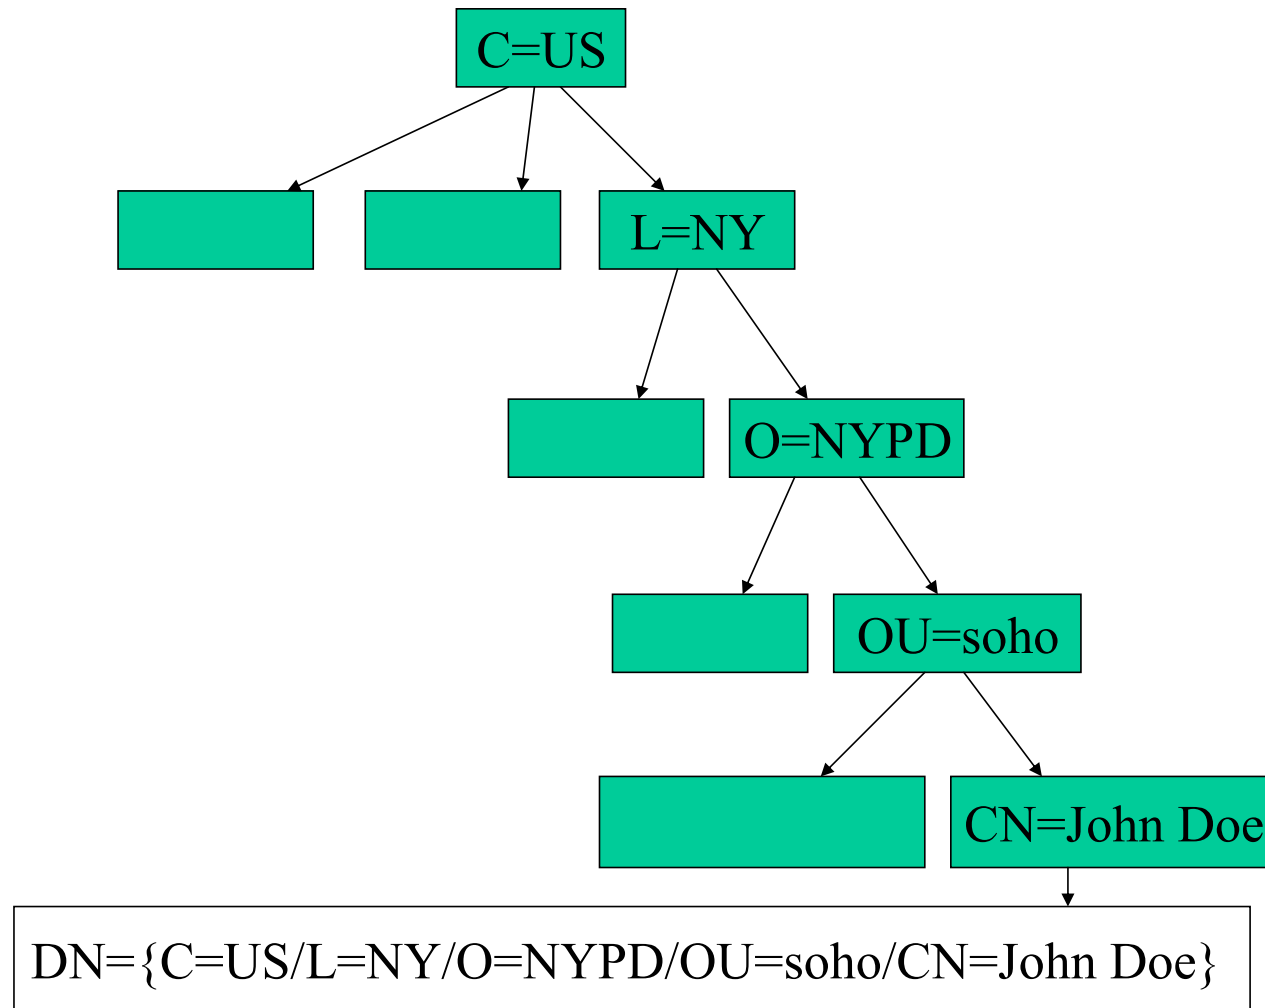


# X.500 Distinguished Names (DN)

- Sequence of keywords, a string value for each of them  
→ *hierarchical DN*.
  - ❑ Keywords facilitate entities sharing same common name.
    - Still uniqueness is not 100% guaranteed.
  - ❑ Meaningful, readable representation.
  - ❑ Distributed directory; each issue manages their issued DNs.

| Keyword | Meaning                |
|---------|------------------------|
| C       | Country                |
| L       | Locality or city name  |
| O       | Organization name      |
| OU      | Organization Unit name |
| CN      | Common Name            |

# Distinguished Name (DN) Hierarchy



---

# Intermediate CAs and Path Verification

---

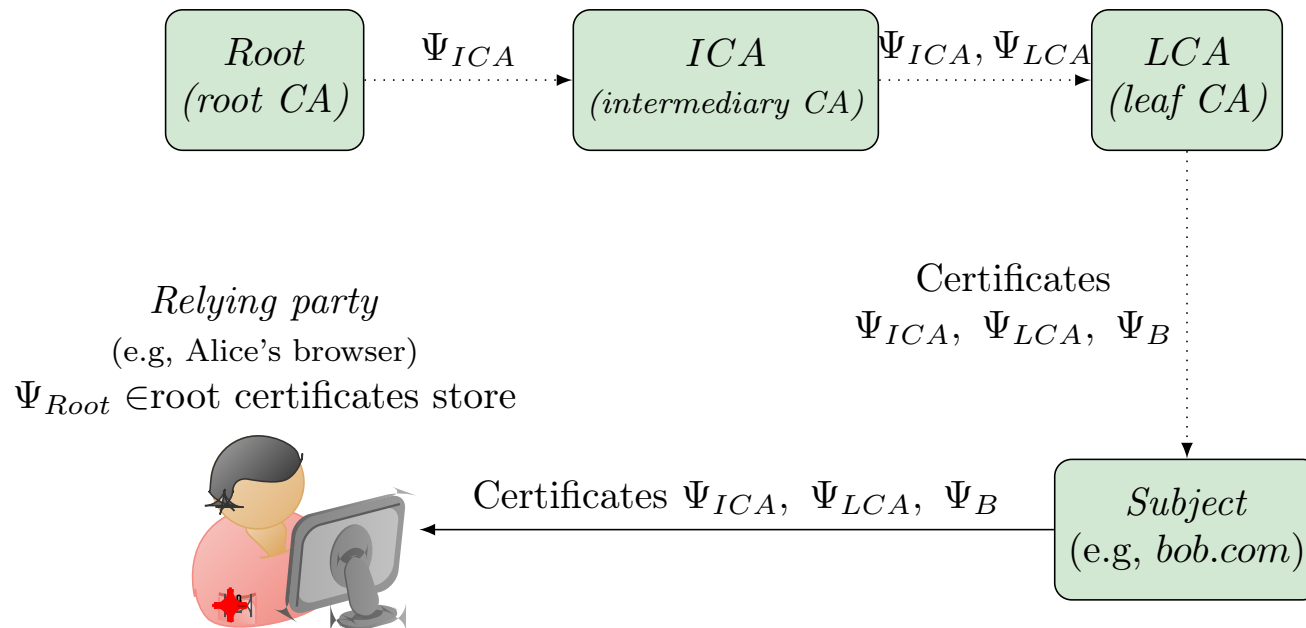
# Why Intermediate CAs?

- Relying parties rely on root CA(s) to establish trust in a certificate of a particular subject party.
  - Large number of subjects to certify.
    - One (or a few) root CAs cannot handle all the load.
  - A root CA certifies other CAs to become intermediate CAs.
    - So the root CA A certifies intermediate CA B, then B will sign certificates for subjects (B is an issuer).
    - Intermediate CAs can certify (beside subjects) other intermediate or leaf CAs.
    - Leaf CAs can certify only subjects.
  - Certificate path validation allows validating such certificates that are issued by intermediate CAs.
    - Like tracing them back to the root CA.
  - *Who certifies a root CA?*
-



# X.509 Validation of Certificate Paths

- Simply, validate all certificates in the chain all the way to the root CA.
- The root CA (self-signed) certificate is in the root store in Alice's browser.
- Let's trace the example below.



---

# Certificate Revocation

# Certificate Revocation

- Reasons for revoking certificates
  - Security issues:
    - Key compromise, CA compromise
  - Administrative issues:
    - Affiliation changed (changing DN or other attribute), public key has been replaced, subject has ceased operation (company dissolving).
- How to inform relying parties? Few options usually under three categories:
  - Prefetch: have revocation info in advance.
  - As-needed: ask for this info when receiving a certificate and want to validate.
  - Neither: does not fall under any of the above, usually called network-assisted techniques.

# Certificate Revocation Techniques

- Prefetch:
  - Cons: higher storage and communication overhead,
  - Pros: lower response delay
- As needed:
  - Cons: higher response delays, reliability issues, privacy concerns.
  - Pros: lower storage and communication overhead
- We will study two techniques:
  - Distribute **Certificate Revocation List (CRL)** -- Prefetch
    - This is part of the X.509 standard.
  - Ask - **Online Certificate Status Protocol (OCSP)** – As needed

# CRLs

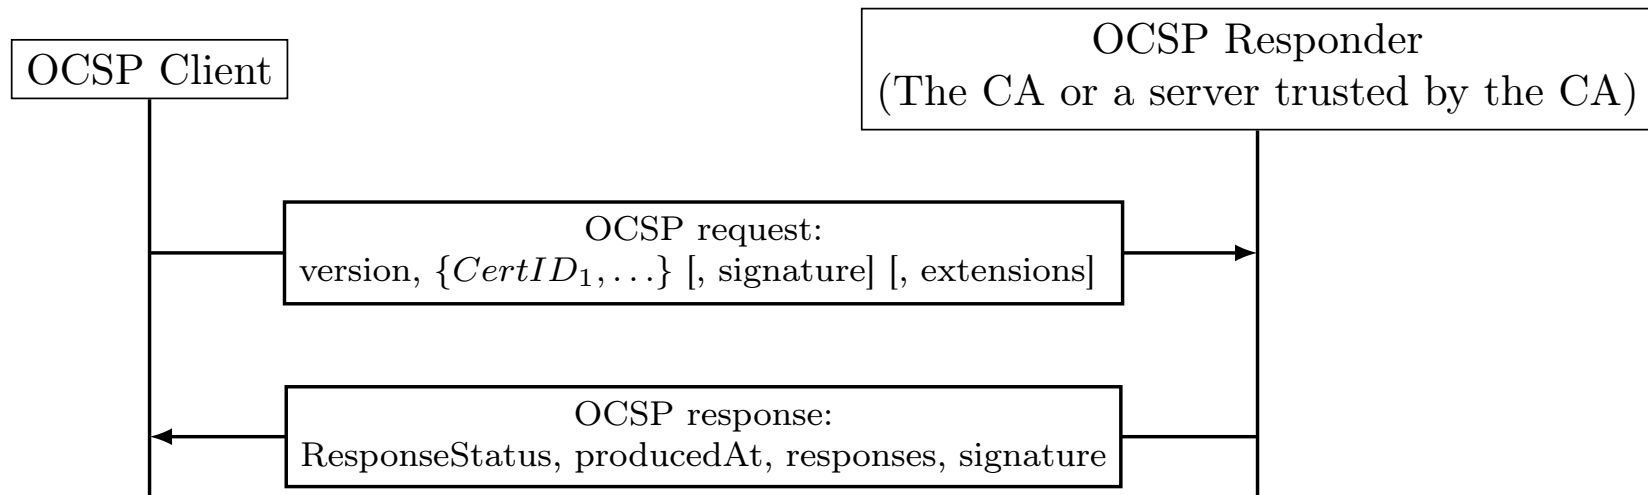
- A certificate revocation list (CRL) is simply a list of revoked certificates.
  - Distributed periodically by CAs.
- If CRLs contain all revoked certificates (which did not expire)... it may be huge!
  - Yes, large storage and communication overhead.
- CRLs are not immediate
  - Who is responsible until CRL is distributed?
  - Frequent CRLs → even more overhead!

# CRLs Optimization Solutions

- More efficient CRL schemes:
  - CRL distribution point: split certificates to several CRLs
  - Authorities Revocation List (ARL): list only revoked CAs
  - Delta CRL – only new revocations since last `base CRL`
    - Need to keep CRLs for long period to check deltas → complicates implementation
- Browsers mostly do not check CRLs. Instead, they usually use:
  - The Online Certificate Status Protocol (OCSP)

# Online Certificate Status Protocol (OCSP)

- Improve efficiency and freshness compared to CRLs.
- Client asks CA about cert during handshake.
- CA signs response (real-time).



---

# OCSP Challenges

- Privacy (expose domain and client to CA), load on CA, response delay, reliability (what if CA fails).
- Ambiguity:
  - When an OCSP server (or CA) cannot resolve the request, it replies with "certificate status is unknown".
- Reliability or failed requests.
  - Client failed to establish a connection with the OCSP server.
  - Or client's request is invalid (not signed, or not authorized), so no response will be received.



---

# Ambiguous/Failed OCSP Responses

- What should the client do?
  - Wait forever – unrealistic!
  - Hard-fail: terminate the connection since certificate is unknown/not received.
    - Safe!
  - Ask user: application display a message asking the user how to proceed.
  - Soft-fail: pretend that a response has been received and continue as the certificate is not revoked.
    - Common choice for browsers!
    - But, a man in the middle (MitM) attacker may block the OCSP response to make a revoked cert go through!

---

# Conclusion

- PKI is an essential component of the Internet.
- Yet, it is a complicated module with many issues related to security, privacy, and performance.
  - To many, this is a solved problem, but that is not the case.
  - Several open questions related to how to detect rogue certificates, how to handle CA failure, revocation, etc., how to audit these parties, how to reduce trust,...
  - How to handle all these issues in an efficient way?
    - Remember, we all want a Web that is highly responsive!

---

# Covered Material From the Textbook

- ❑ Chapter 8:
  - ❑ Sections 8.1,
  - ❑ and Sections 8.3 and 8.4 (only the topics we covered from both sections)

---

# Thank You!

